

Error and Flow Control in TCP

Dr. Manas Khatua
Assistant Professor
Dept. of CSE
IIT Jodhpur

E-mail: manaskhatua@iitj.ac.in

SEQ and ACK numbers in TCP

- TCP views data as an unstructured, but ordered, **stream of bytes**
- The **SEQ number for a segment** is the byte-stream number of the **first byte** in the segment.
- The **ACK number** that Host A puts in its segment is the **SEQ number of the next byte** Host A is expecting from Host B.

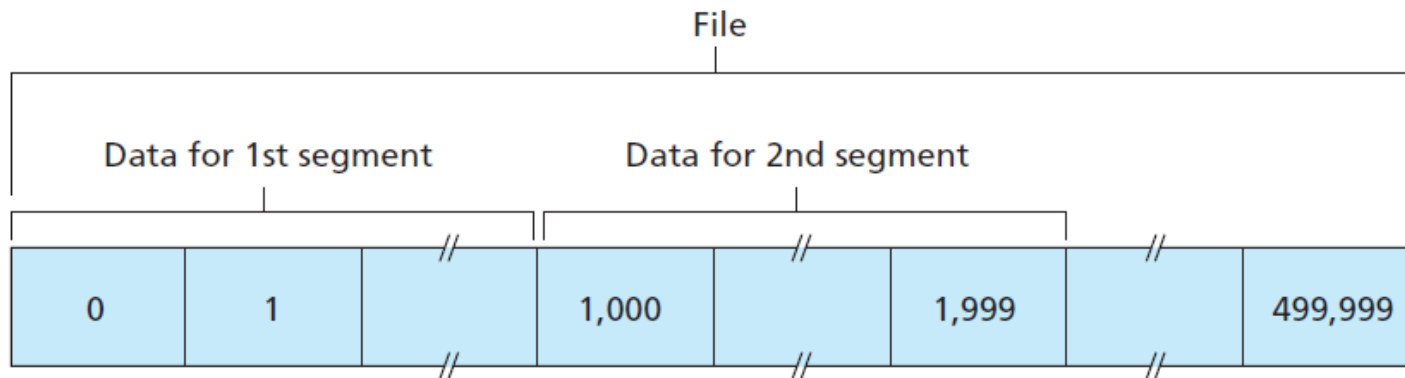


Figure 3.30 ♦ Dividing file data into TCP segments

Error Control in TCP

- Error control in TCP is done by :
 - checksum, ACK, time-out
- By default TCP uses cumulative ACK
- When does a receiver generate ACK?
 - Rule-1: when node A sends data to node B, it piggybacks ACK
 - Rule-2: the receiver has no data to send and it receives an in-order segment, it delays sending ACK
 - Rule-3: there should not be more than two in-order unACKed segments at any time (it is delayed ACK)
 - Rule-4: when a segment arrives with an out-of-order sequence number; or, it is fast retransmission of missing segments
 - Rule-5: when a missing segment arrives
 - Rule-6: If a duplicate segment arrives

Cont...



- When retransmission happens?
 - **After time-out**: sending TCP maintains one retransmission time-out (RTO) for each connection
 - **Three duplicate ACK rule**: if three duplicate ACK (i.e., an original ACK + three exactly identical copies) arrive for a segment, the next segment is retransmitted without waiting for the time-out.
 - Why **three duplicate ACK**?
 - Since TCP does not know whether a duplicate ACK is caused by a lost segment or just a reordering of segments, it waits for at max two duplicate ACK. If three or more duplicate ACKs are received in a row, it is a strong indication that a segment has been lost.
- Data may arrive **out of order** and be temporarily stored by the receiving TCP, but **TCP guarantees** that no out-of-order data are delivered to the process.

RTT Estimation and Timeout

- TCP uses a timeout/retransmit mechanism to recover from lost segments.
- The **timeout** should be larger than the connection's round-trip time (**RTT**)
- **Few Questions:**
 - How much larger? How should the RTT be estimated in the first place?
 - Should a timer be associated with each and every unACKed segment?
- The sample RTT, denoted *SampleRTT*, for a segment is
 - the amount of time between when the segment is sent (that is, passed to IP) and when an ACK for the segment is received.
- But, the *SampleRTT* values will fluctuate from segment to segment due to congestion in the routers and to the varying load on the end systems.
- **Solution:**
 - TCP maintains an average, called *EstimatedRTT*, of the *SampleRTT* values.
 - Exponentially weighted moving average (EWMA)

$$\text{EstimatedRTT} = (1 - \alpha) \times \text{EstimatedRTT} + \alpha \times \text{SampleRTT}$$

The recommended value of α is $= 1/8$

Cont...



- It is also valuable to have a measure of the **variability of the RTT**.

$$\text{DevRTT} = (1 - \beta) \times \text{DevRTT} + \beta \times |\text{SampleRTT} - \text{EstimatedRTT}|$$

The recommended value of β is $1/4$.

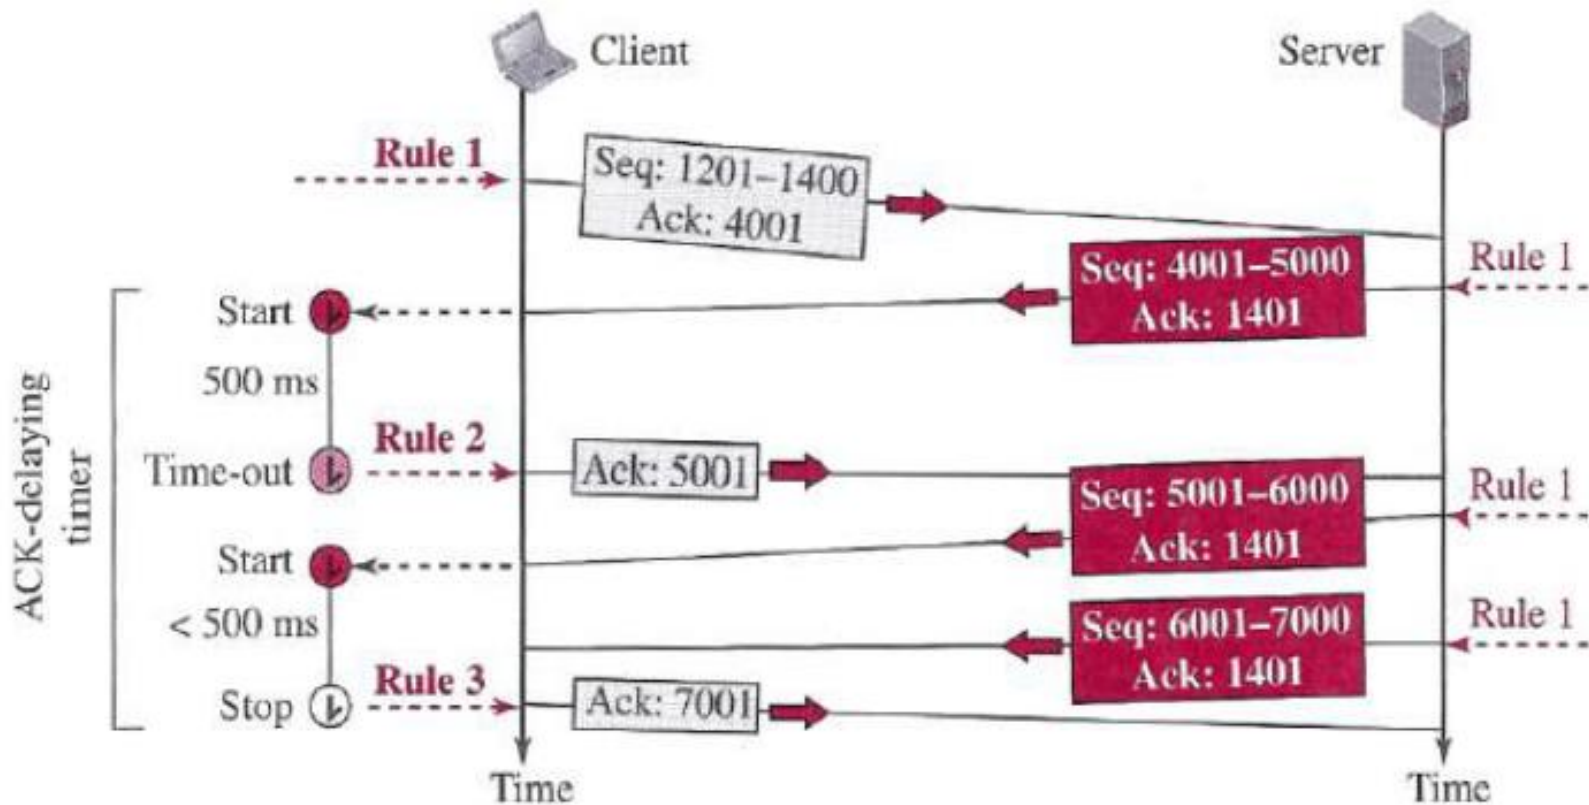
- For little fluctuation, the DevRTT will be small; on the other hand, for lot of fluctuation, DevRTT will be large.
- It is desirable to set the timeout equal to the *EstimatedRTT* plus some margin.
- The margin should be large when there is a lot of fluctuation in the *SampleRTT* values; it should be small when there is little fluctuation.

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \times \text{DevRTT}$$

- Another Question:** How long the receiver waits before sending a stand-alone ACK to acknowledge data on a TCP socket?
 - Delayed ACK* was invented to reduce the number of ACKs required to acknowledge the segments, so protocol overhead is reduced.
 - A host may delay sending an ACK response by up to **500 ms**.
 - However, a stand-alone ACK is sent if two full packets worth of data arrive before the delayed ACK timer expires.

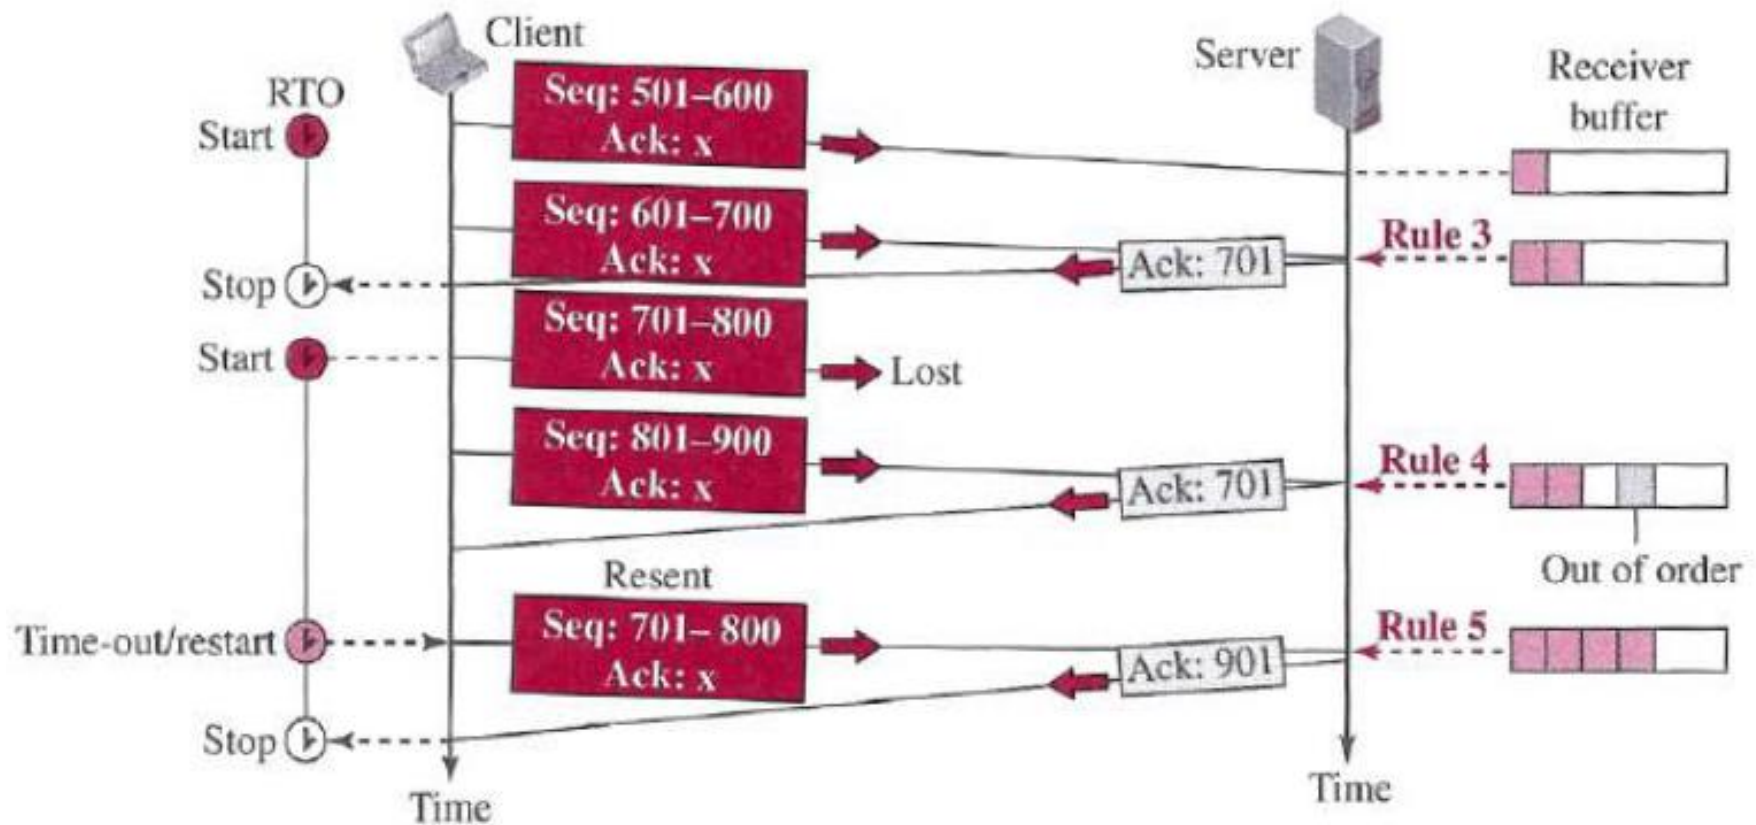
Example : Normal Scenario

- The client TCP sends one segment (2000 byte);
- server TCP sends three segments (3 x 1000 byte).

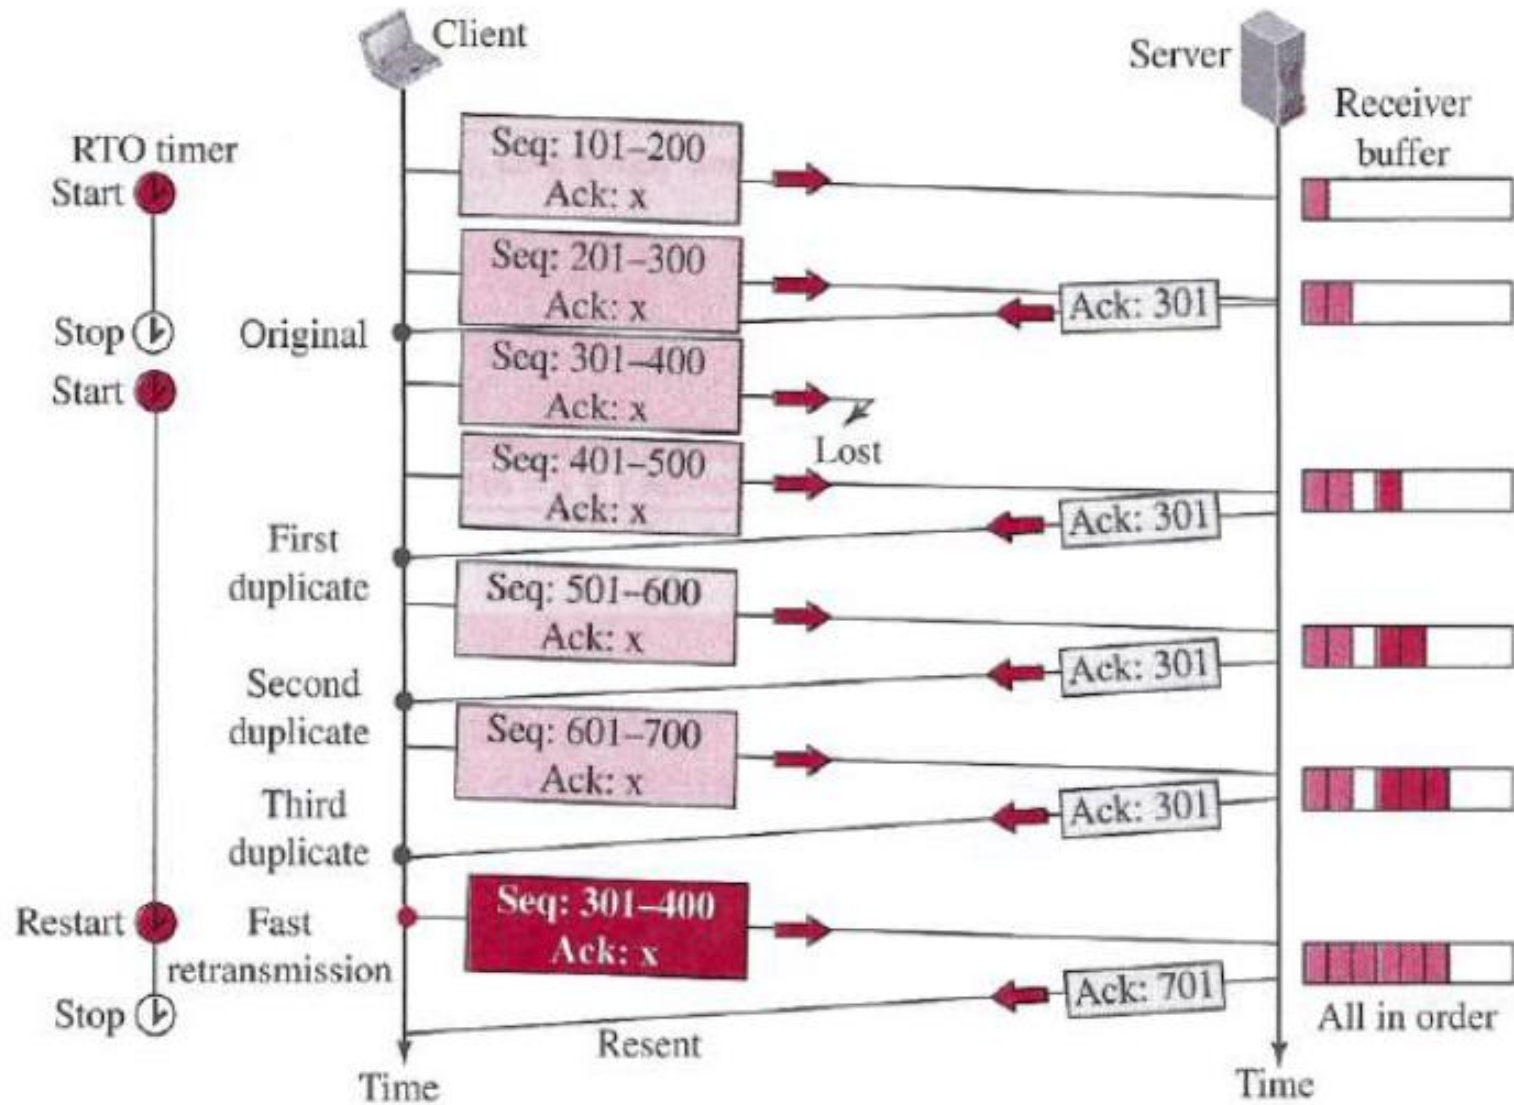


Example : Lost Segment Scenario

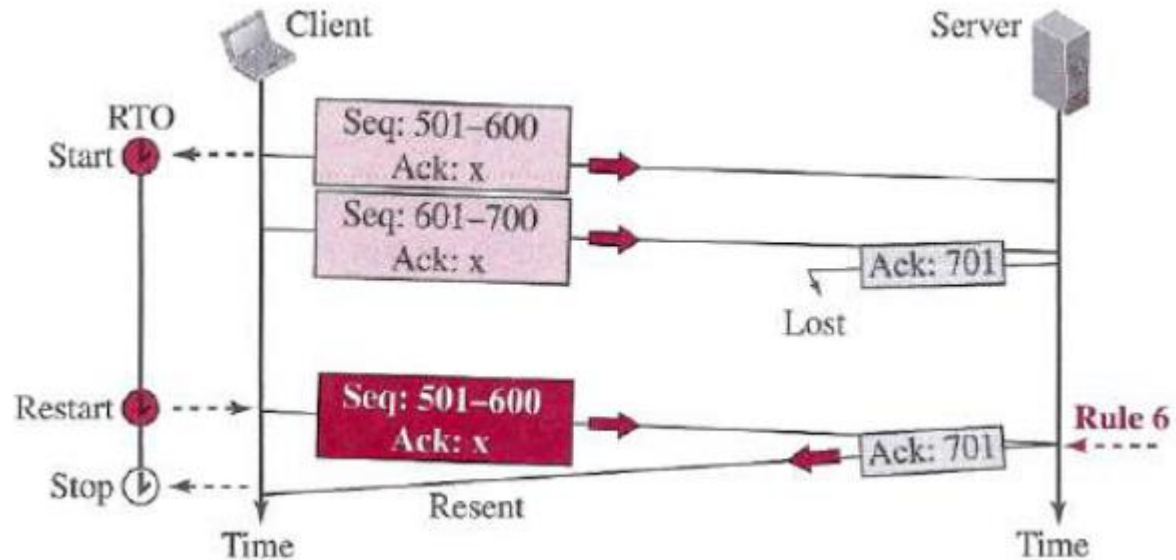
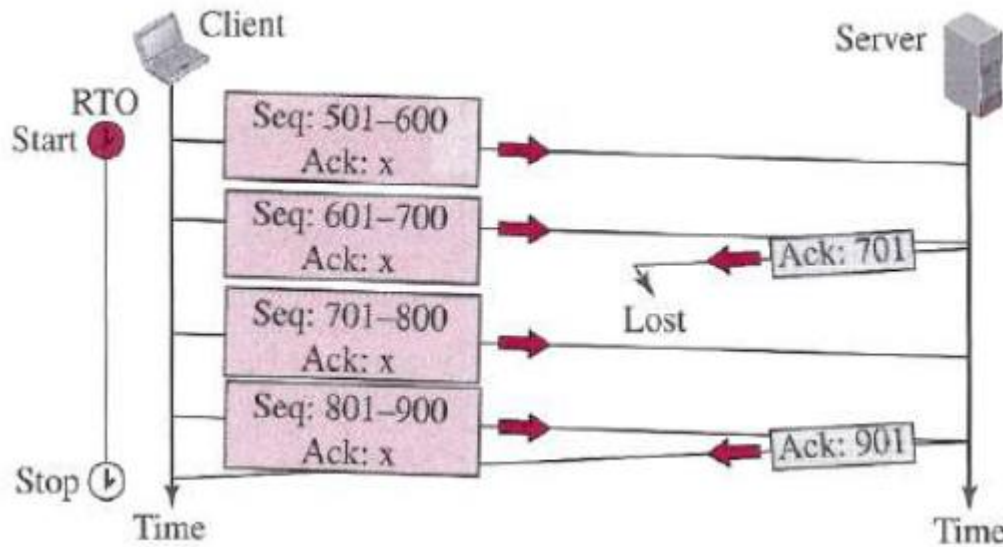
- assuming that data transfer is unidirectional
– from client to server



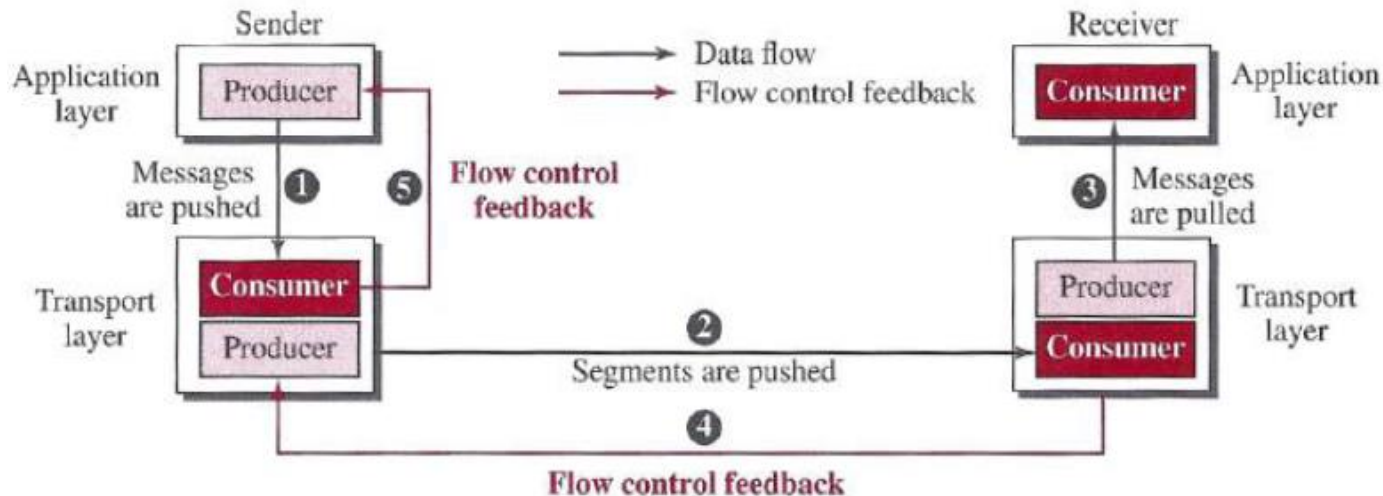
Example : Fast Retransmission



Example : Lost ACK



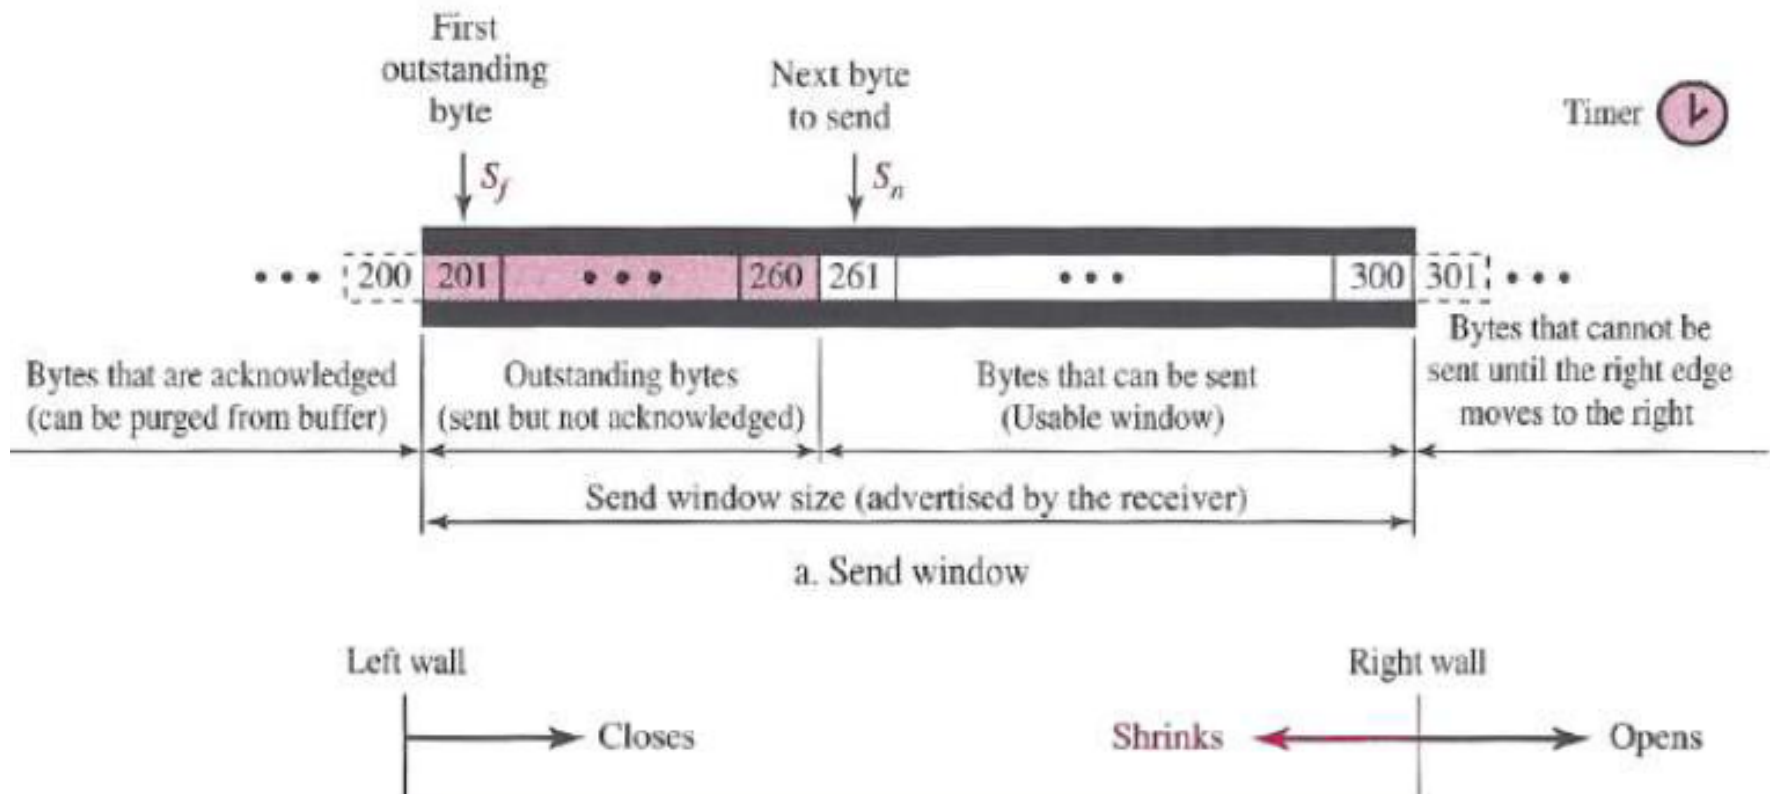
Flow Control in TCP



- The **receiving TCP** controls the **sending TCP**; the **sending TCP** controls the **sending process**.
- No flow control between receiving TCP and receiving process.
- To **achieve flow control**:
 - TCP forces the sender and the receiver to **adjust their window sizes**, although the size of the buffer for both parties is fixed when the connection is established.
 - The **opening, closing, and shrinking** of the send window is **controlled by the receiver**.

Send Window in TCP

- TCP uses Send window & Receive window
- Let send window size = 100

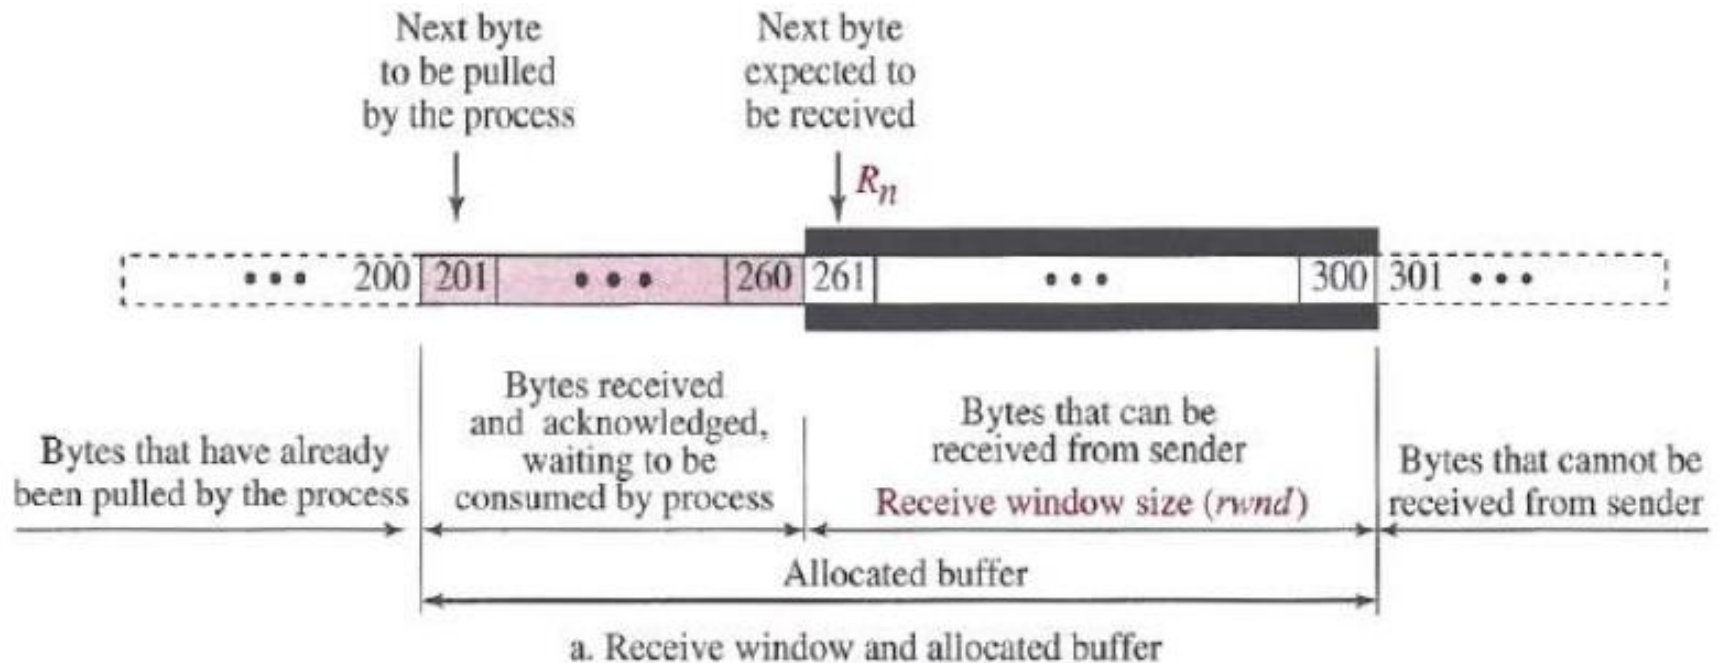


Modified SR (for Send Window)



- Sending window in TCP follows **Selective-Repeat (SR)** protocol with few modifications
 - The window size in SR is the number of packets, but the window size in TCP is the **number of bytes**.
 - TCP can **store data** received from the process and send them later
 - SR requires individual ACK of each packet that was sent; but **TCP sends ACK for the next packet that it is expecting** (like Go-back-N)
 - SR protocol may use several timers for general transmission and selected retransmission, but TCP protocol uses only **one timer**.
 - **Window size can be changed dynamically** in TCP

Receive Window in TCP

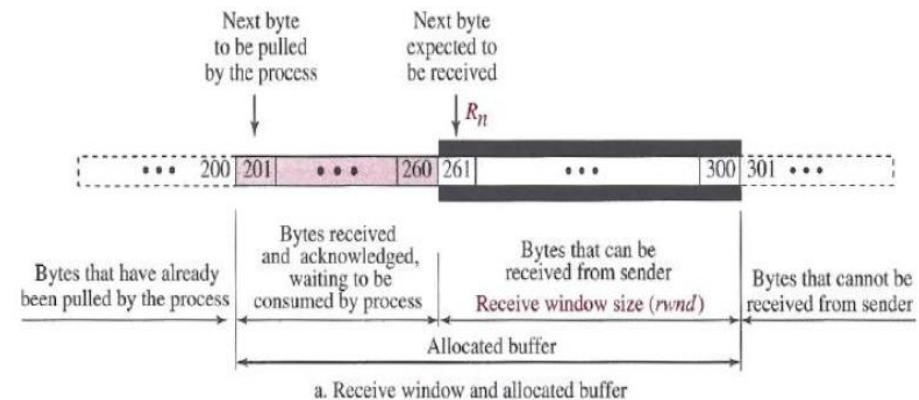



Modified SR (for Receive Window)

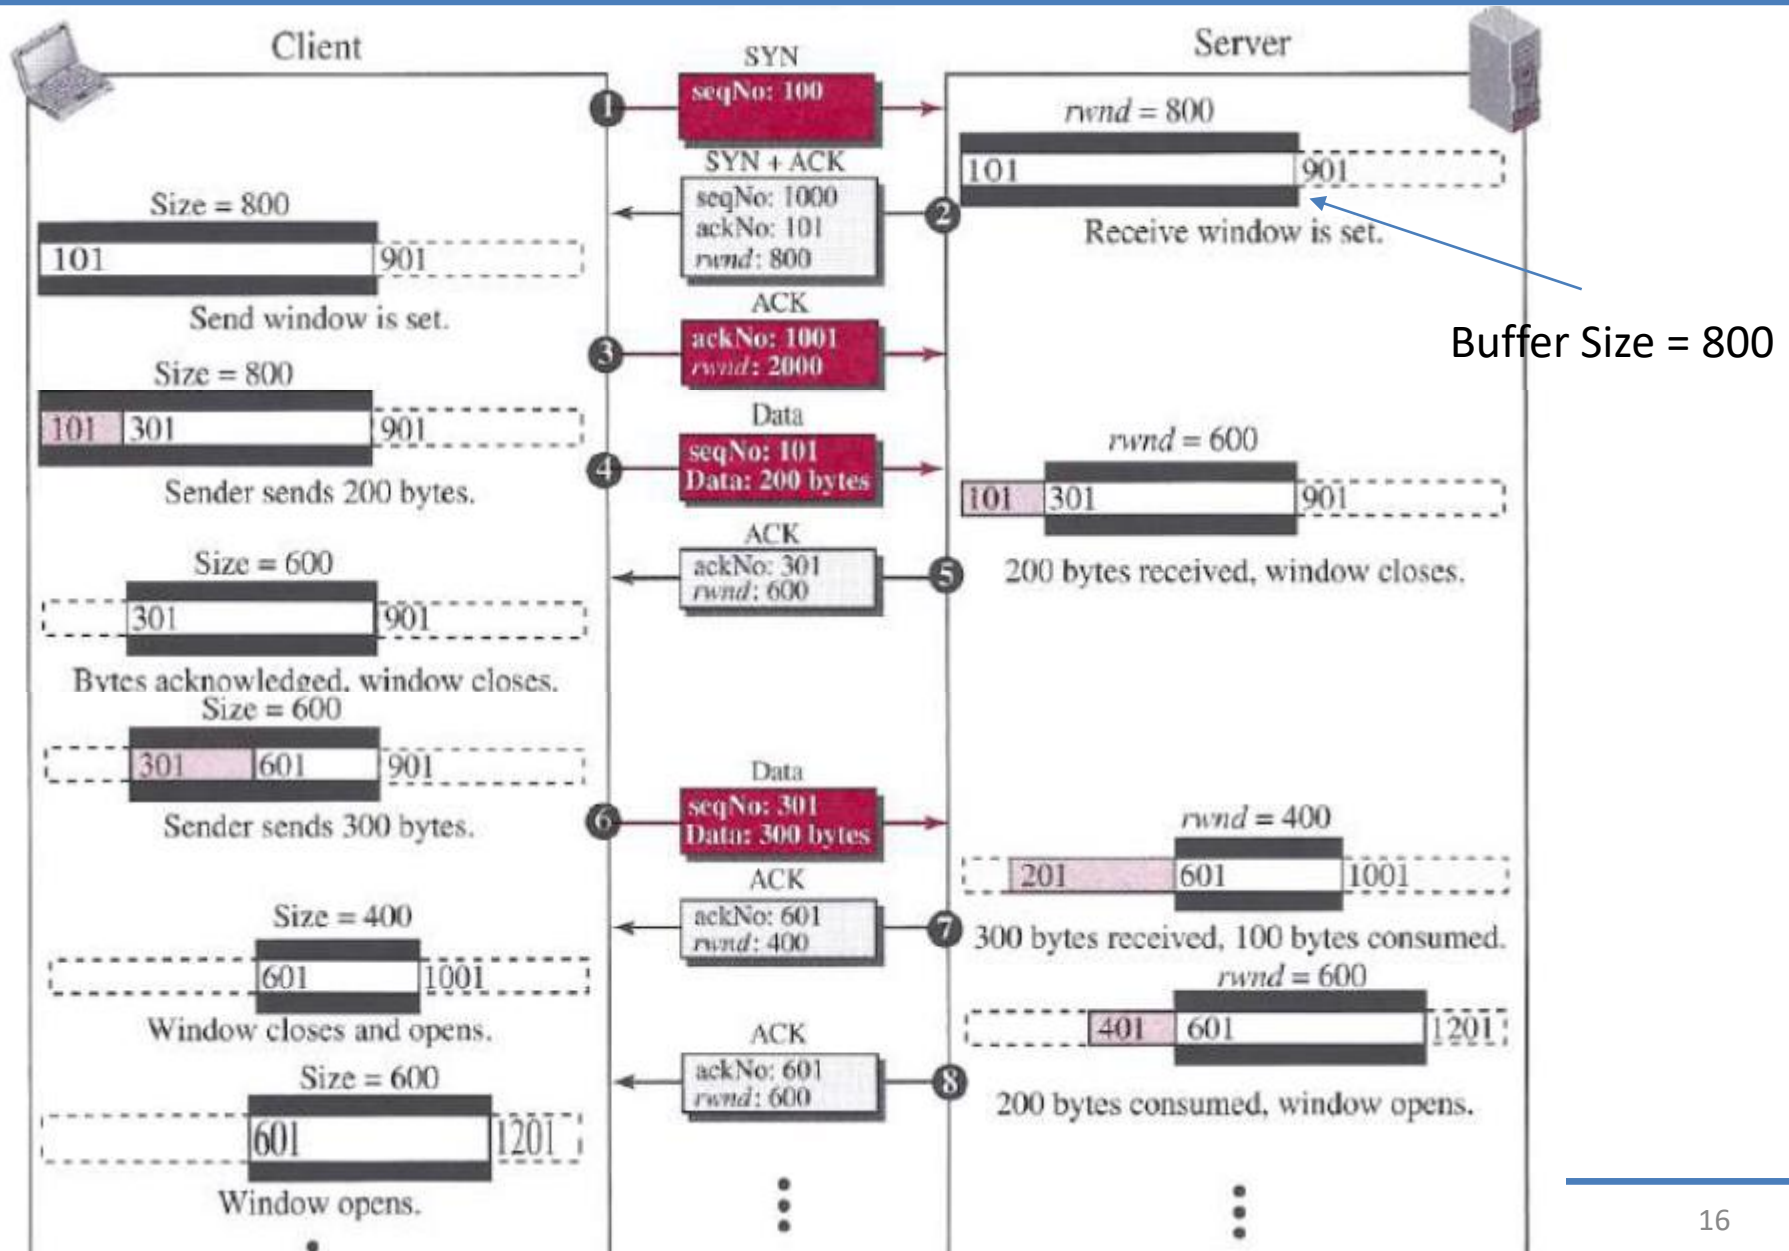
- Receive window in TCP is little different than that in SR
 - TCP allows the receiving process to **pull data** at its own pace.
 - The receive window size (***rwnd***) determines the number of bytes that the receive window can accept from the sender before being overwhelmed (**flow control**).

$$rwnd = \text{buffer size} - \text{number of bytes waiting to be pulled}$$

- **ACK** in SR is selective, but ACK in TCP is **cumulative**.
- Similarly, retransmission is selective in SR, but **oldest unACKed** segment is retransmitted in TCP




Example (oneway from client to server)



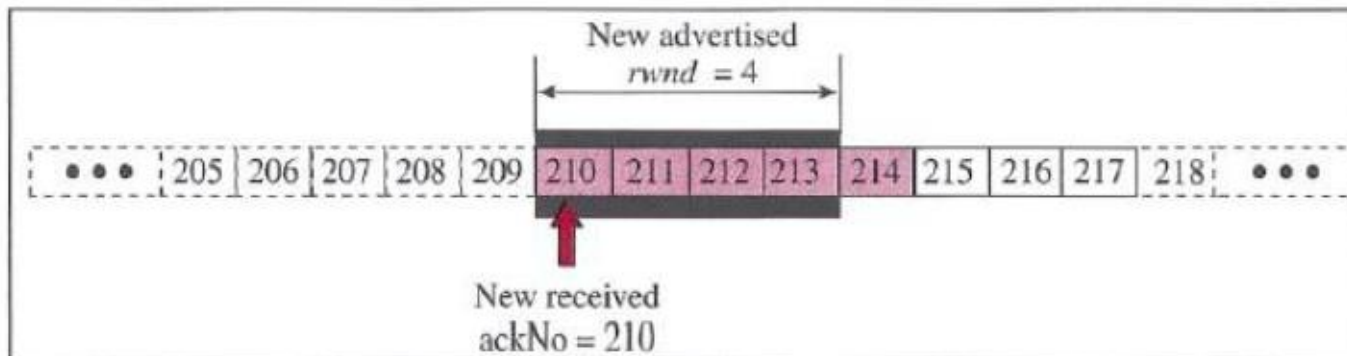
Shrinking of Windows

$\text{new ackNo} + \text{new } rwnd \geq \text{last ackNo} + \text{last } rwnd$

If $rwnd=0$, it instructs for “window shutdown”



a. The window after the last advertisement



b. The window after the new advertisement; window has shrunk

Silly Window Syndrome



- Performance issue occurs when
 - Sending **application** program **creates data slowly**
 - OR, Receiving **application** program **receives data slowly**
 - OR, For **both the above**
- Example:
 - Sending process generating each byte very slowly
 - Sending TCP sends many 41 bytes segment (20 byte TCP header + 20 byte IP header + 1 byte data)
- Two types to address
 - Syndrome created by sender
 - Syndrome created by receiver

Solution



- **Naïve solution** faces a trade-off optimization
 - If TCP waits too long, it may delay the process
 - If TCP does not wait for long, it may end up sending small segment
- **Better Solution for sender:** Nagle's Algorithm
 - Sending TCP sends the **1st segment** as it is
 - **2nd segment onwards**, the sending TCP accumulates data in sending buffer and waits until
 - Either the receiving TCP sends an ACK
 - Or enough data have accumulated **to fill the maximum-size segment**
- **Better Solution for receiver:** Clark's two algorithms
- First algorithm:
 - send an ACK as soon as the data arrive,
 - but to announce a window size of zero until
 - either there is enough space to accommodate a segment of maximum size
 - or until at least half of the receive buffer is empty.
- Second algorithm:
 - delay sending the ACK.
 - The receiver waits until there is a decent amount of space in its incoming buffer before ACKing the arrived segments.

Thanks!