

# CS578: Internet of Things

# MQTT: Message Queuing Telemetry Transport



# Dr. Manas Khatua

Assistant Professor

Dept. of CSE, IIT Guwahati

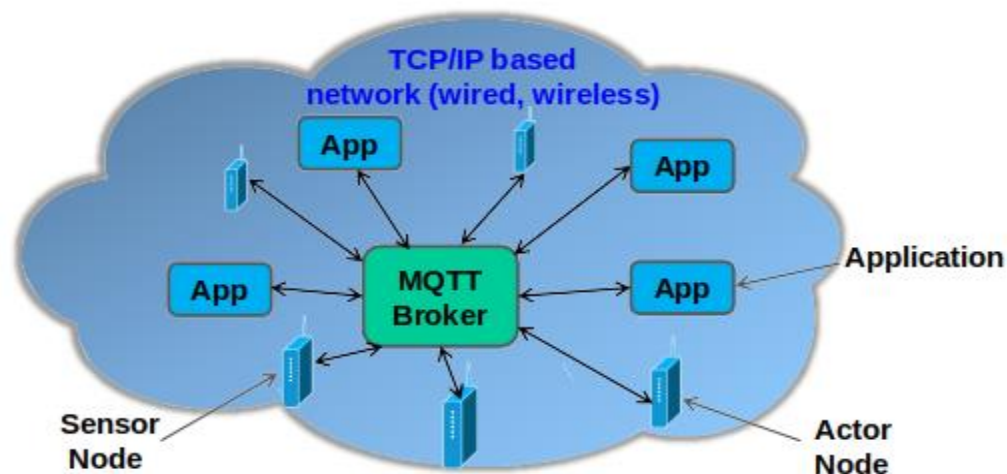
E-mail: [manaskhatua@iitg.ac.in](mailto:manaskhatua@iitg.ac.in)

# What is MQTT?

- MQTT is message queueing and transport protocol.
- **Lightweight protocol**
- Suited for the **transport of telemetry data** (sensor and actor data), and
- thus for **M2M** (Mobile to Mobile), **WSN** (Wireless Sensor Networks) application

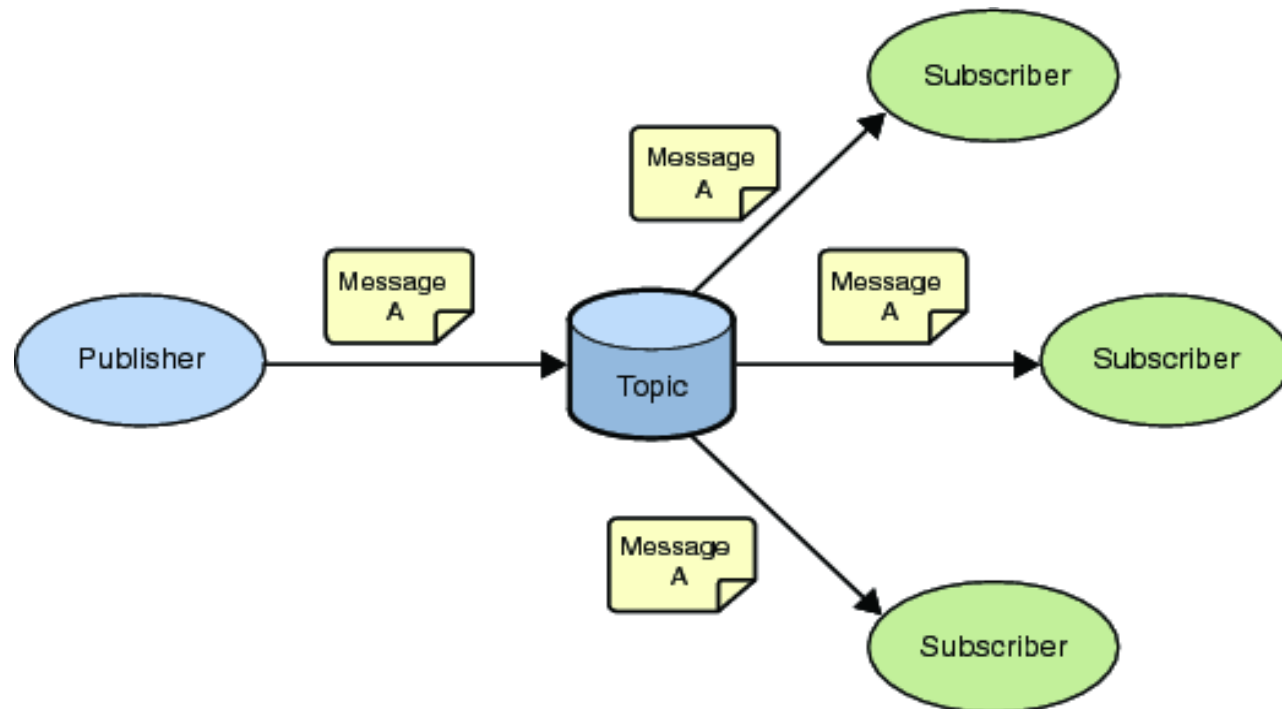
## Example:

- **Light sensor** continuously sends sensor data to the **broker**.
  - **Building control application** receives sensor data from the broker and decides to activate the **blinds**.
  - **Application** sends a blind activation message to the **blind actor** node through the **broker**.
- 
- Invented by Dr. Andy Stanford-Clark of IBM and Arlen Nipper of Arcom (now Eurotech) in 1999
  - Used by
    - Amazon Web Services (AWS),
    - IBM WebSphere MQ,
    - Microsoft Azure IoT,
    - Adafruit,
    - Facebook Messenger,
    - etc.



# MQTT Characteristics

- **Asynchronous communication model** with messages (events)
- Low overhead (**2 bytes header**) for low network bandwidth applications
- **Publish / Subscribe** (PubSub) model
- **Decoupling** of data producer (publisher) and data consumer (subscriber) through topics (message queues)
- Runs on connection-oriented transport (**TCP**). To be used in conjunction with **6LoWPAN** (TCP header compression)
- MQTT caters for (wireless) **network disruptions**



# Publish Subscribe Messaging



## Terminology

- A producer sends (**publishes**) a message (**publication**) on a topic (**subject**)
- A consumer **subscribes** (makes a subscription) for messages on a topic (**subject**)
- A message **server / broker** matches publications to subscriptions

## Who will get the message ?

- If **no matches** the message is **discarded**
- If **one or more matches** the message is delivered to **each** matching subscriber/consumer

## Topic

- A topic forms the namespace is hierarchical with each “**sub topic**” separated by a /
- An **example** topic space :
  - **A house publishes information about itself on:**  
`<country>/<region>/<town>/<postcode>/<house>/energyConsumption`  
`<country>/<region>/<town>/<postcode>/<house>/solarEnergy`
  - **And subscribes for control commands:**  
`<country>/<region>/<town>/<postcode>/<house>/thermostat/setTemp`

# Publish Subscribe Messaging



## Wildcards

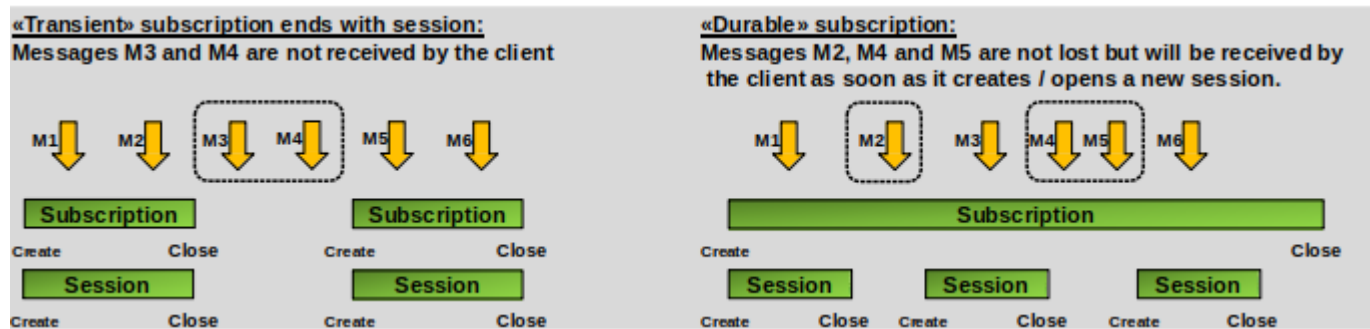
- A subscriber can subscribe to an absolute topic or can use **wildcards**:
  - Single-level wildcards “+” can appear anywhere in the topic string  
**For example:**
    - Energy consumption for 1 house in Hursley
      - *UK/Hants/Hursley/SO212JN/1/energyConsumption*
    - Energy consumption for all houses in Hursley
      - *UK/Hants/Hursley/+/+/energyConsumption*
  - Multi-level wildcards “#” must appear at the end of the string  
**For example:**
    - Details of energy consumption, solar and alarm for all houses in SO212JN
      - *UK/Hants/Hursley/SO212JN/#*

## NOTE :

- Wildcards must be **next to a separator**
- **Cannot** be used wildcards when **publishing**

# Publish Subscribe Messaging

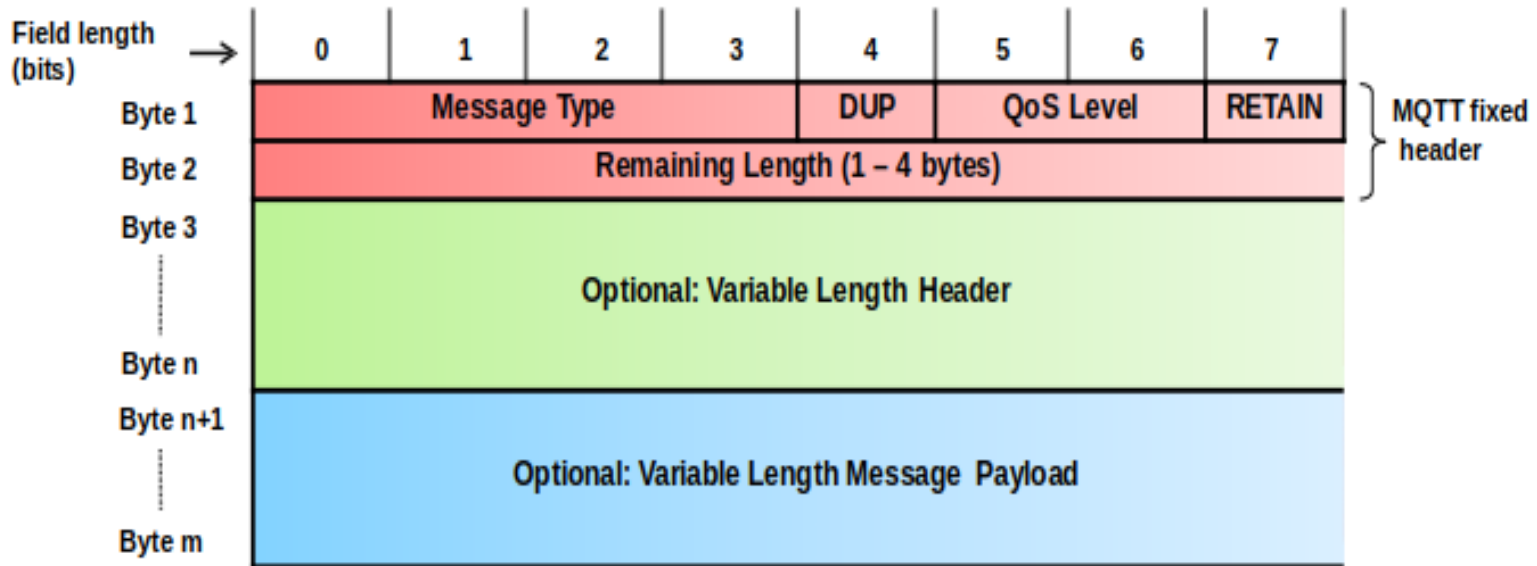
- A subscription can be durable or non durable
- Durable:
  - Once a **subscription** is in place a **broker** will forward matching messages to the **subscriber**:
    - Immediately if the subscriber is connected.
    - If the subscriber is not connected messages are stored on the server/broker until the next time the subscriber connects.



- Non-durable(Transient):
  - The subscription **lifetime** is the **same** as the time the subscriber is **connected** to the server / broker

# MQTT message format

- MQTT messages contain a mandatory **fixed-length header (2 bytes)** and an optional **message-specific variable length header** and **message payload**.
- MQTT uses **network byte and bit ordering**.



# Overview of fixed header fields



Message fixed header field	Description / Values	
Message Type	0: Reserved	8: SUBSCRIBE
	1: CONNECT	9: SUBACK
	2: CONNACK	10: UNSUBSCRIBE
	3: PUBLISH	11: UNSUBACK
	4: PUBACK	12: PINGREQ
	5: PUBREC	13: PINGRESP
	6: PUBREL	14: DISCONNECT
	7: PUBCOMP	15: Reserved
DUP	Duplicate message flag. Indicates to the receiver that this message may have already been received. 1: Client or server (broker) re-delivers a PUBLISH, PUBREL, SUBSCRIBE or UNSUBSCRIBE message (duplicate message).	
QoS Level	Indicates the level of delivery assurance of a PUBLISH message. 0: At-most-once delivery, no guarantees, «Fire and Forget». 1: At-least-once delivery, acknowledged delivery. 2: Exactly-once delivery.	
RETAIN	1: Instructs the server to retain the last received PUBLISH message and deliver it as a first message to new subscriptions.	
Remaining Length	Indicates the number of remaining bytes in the message, i.e. the length of the (optional) variable length header and (optional) payload.	



# RETAIN (keep last message)

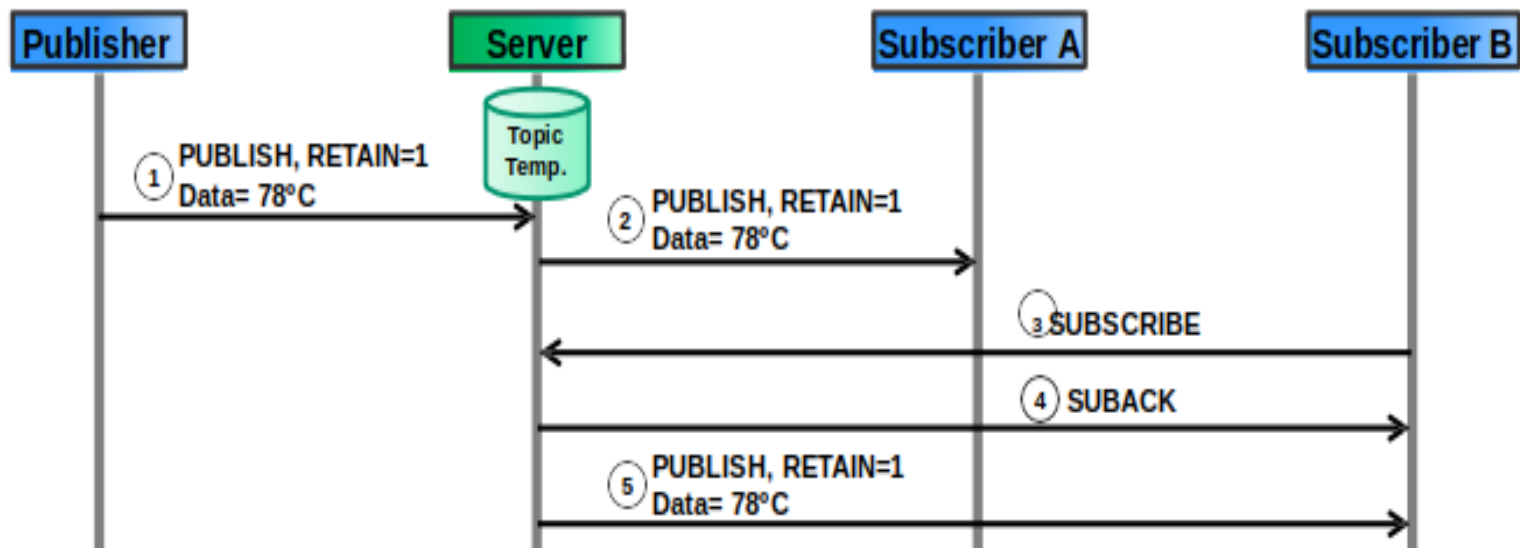
**RETAIN=1** in a **PUBLISH** message instructs the server to keep the message for this topic. When a new client subscribes to the topic, the server sends the retained message.

- **Typical application scenarios:**

- Clients publish only **changes** in data, so subscribers receive the last known good value.

- **Example:**

- Subscribers receive **last known temperature value** from the temperature data topic.
- **RETAIN=1** indicates to subscriber B that the message may have been published some time ago.



# Remaining length (RL)




The remaining length field encodes the sum of the lengths of:

1. **(Optional)** variable length header
2. **(Optional)** payload

To save bits, remaining length is a **variable length field** with **1...4 bytes**.

The **most significant bit** of a length field byte has the meaning «continuation bit» (CB). If more bytes follow, it is set to 1.

Remaining length is encoded as  $a * 128^0 + b * 128^1 + c * 128^2 + d * 128^3$  and placed into the RL field bytes as follows:



Key:  
LSB: Least Significant Byte  
MSB: Most Significant Byte


**Example 1:**  $RL = 364 = 108 * 128^0 + 2 * 128^1 \rightarrow a=108, CB0=1, b=2, CB1=0, c=0, d=0, CB2=0$

**Example 2:**  $RL = 25'897 = 41 * 128^0 + 74 * 128^1 + 1 * 128^2 \rightarrow a=41, CB0=1, b=74, CB1=1, c=1, CB2=0, d=0$

# MQTT QoS



- MQTT provides the typical delivery quality of service (QoS) levels of message oriented middleware.
- Even though TCP/IP provides guaranteed data delivery, **data loss** can still occur if a **TCP connection breaks down** and messages in transit are **lost**.
- Therefore MQTT adds **3 quality of service** levels on top of TCP



## QoS level 0:

- **At-most-once delivery** («best effort»).
- Messages are delivered according to the delivery **guarantees of the underlying network** (TCP/IP).
- **Example application:** Temperature sensor data which is regularly published. **Loss of an individual value** is not critical since applications (consumers of the data) will anyway **integrate the values over time** and loss of individual samples is not relevant.

## QoS level 1:

- **At-least-once delivery.**
- Messages are guaranteed to arrive, but there may be **duplicates**.
- **Example application:** A door sensor senses the door state. It is important that door state changes (closed->open, open->closed) are published **loss-lessly** to subscribers (e.g. alarming function). Applications simply discard duplicate messages by evaluating the message ID field.

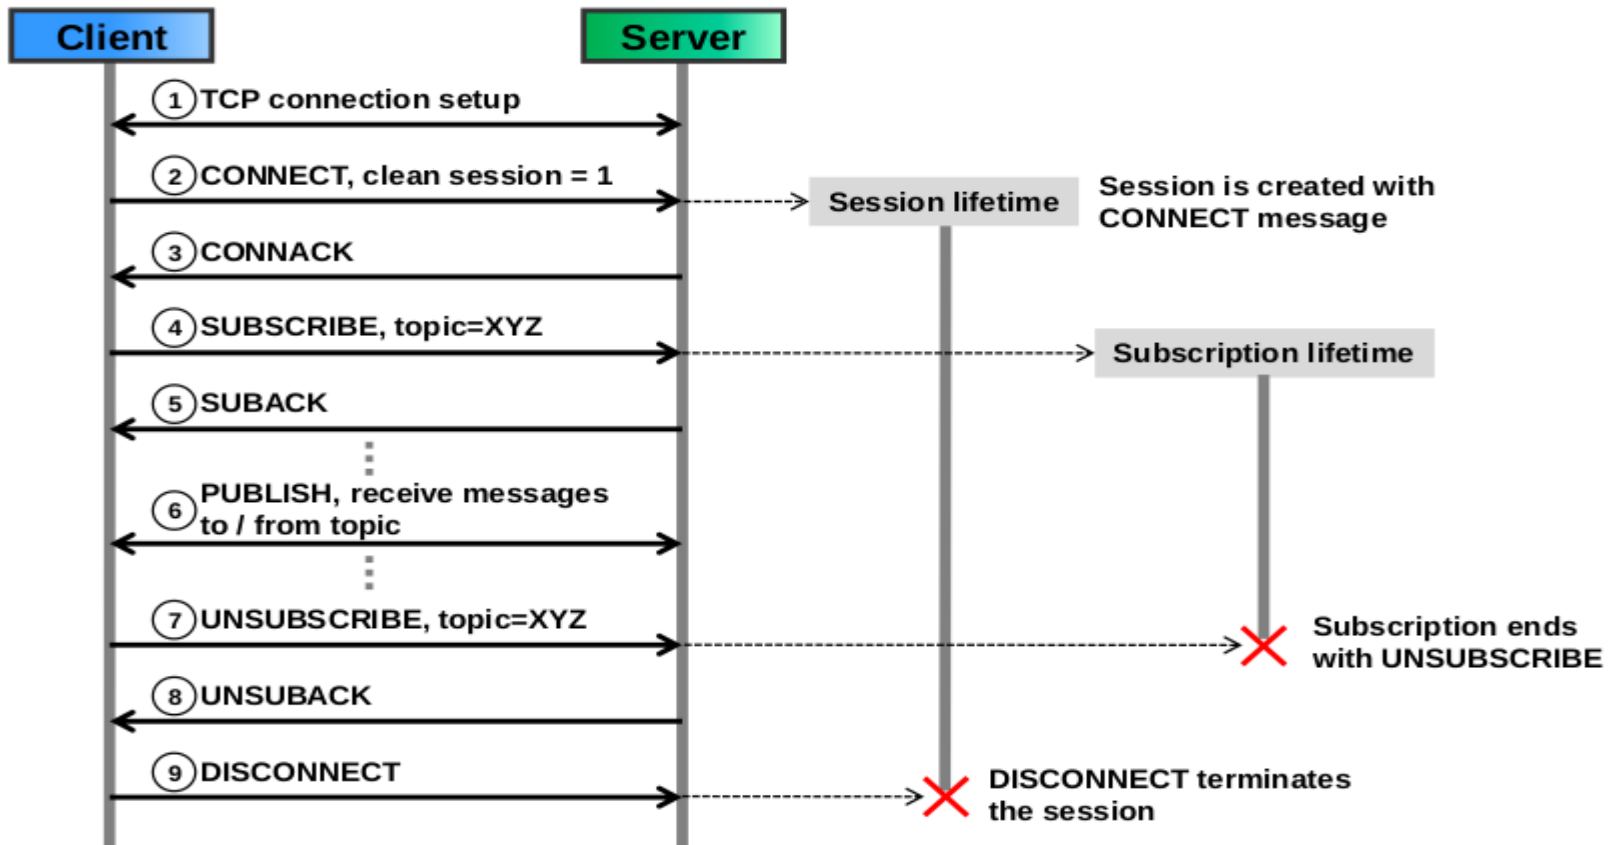
## QoS level 2:

- **Exactly-once delivery.**
- This is the highest level that also incurs **most overhead** in terms of control messages and the need for **locally storing** the messages.
- **Exactly-once** is a **combination** of **at-least-once** and at-**most-once** delivery guarantee.
- **Example application:** Applications where duplicate events could lead to incorrect actions, e.g. sounding an alarm as a reaction to an event received by a message.

# CONNECT and SUBSCRIBE msg flow

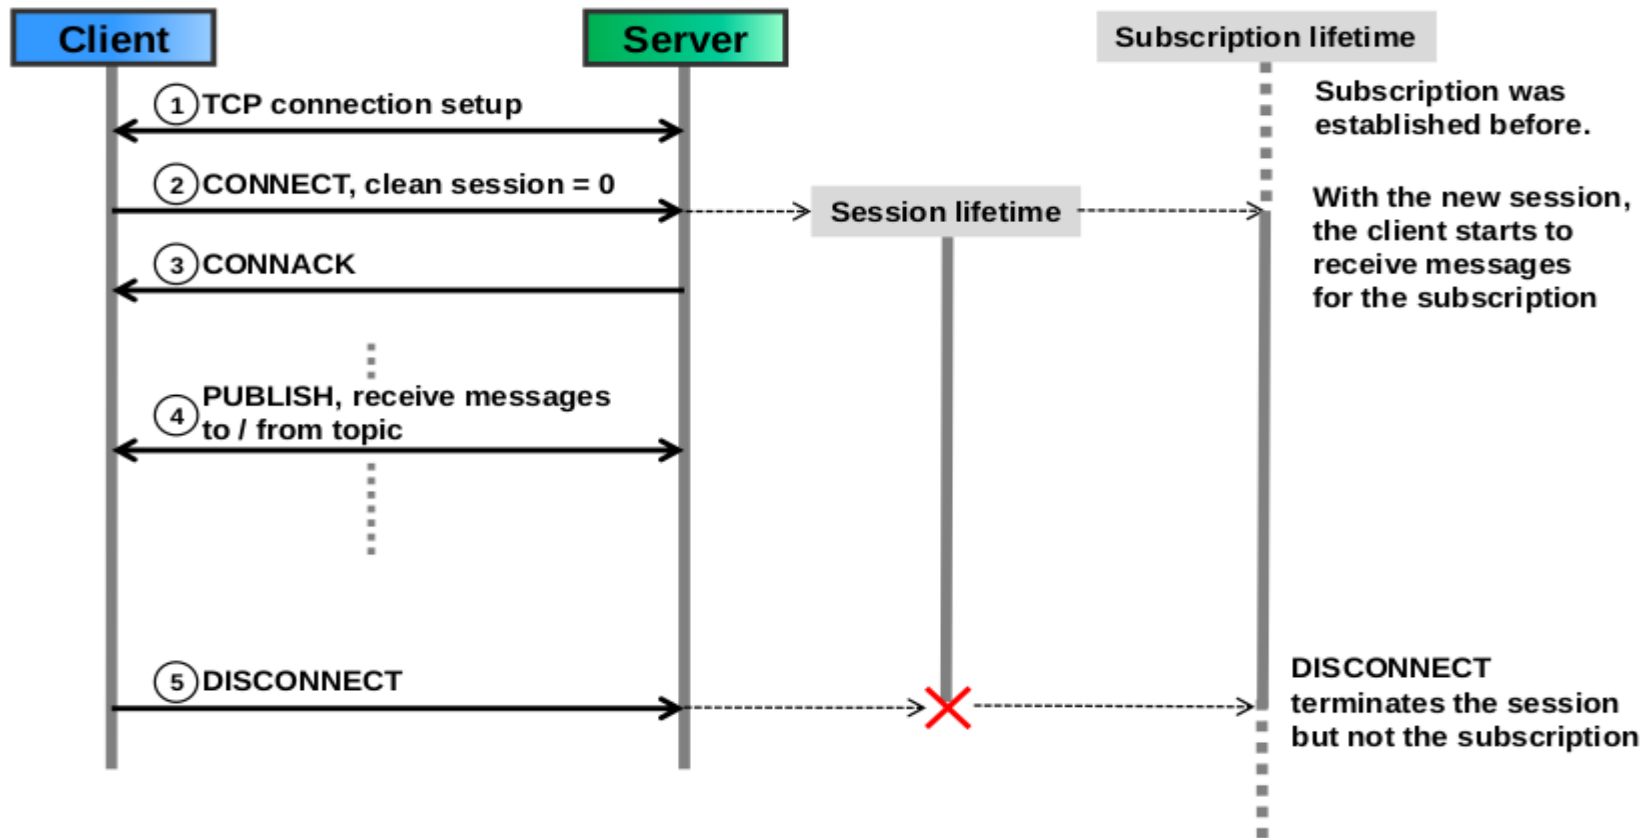
## Case 1:

- Session and subscription setup with clean session flag = 1 («transient» subscription)



## Case 2:


- Session and subscription setup with clean session flag = 0 («durable» subscription)



# PUBLISH msg flows

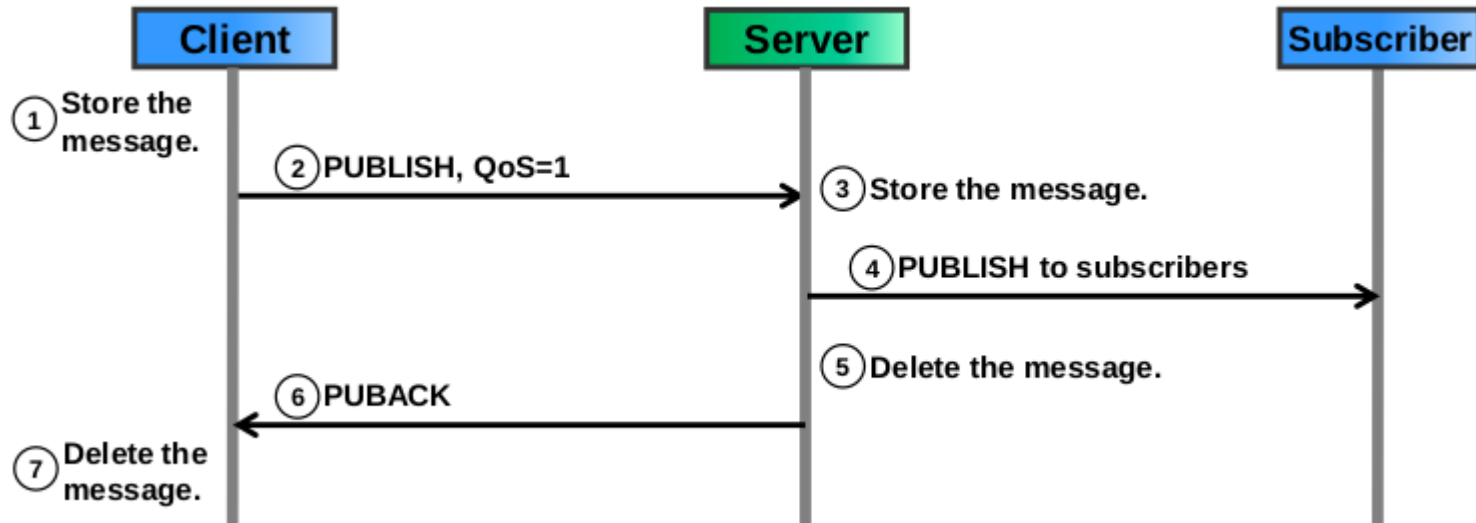
## QoS level 0:

- With QoS level 0, a message is delivered with **at-most-once delivery semantics** («fire-and-forget»).



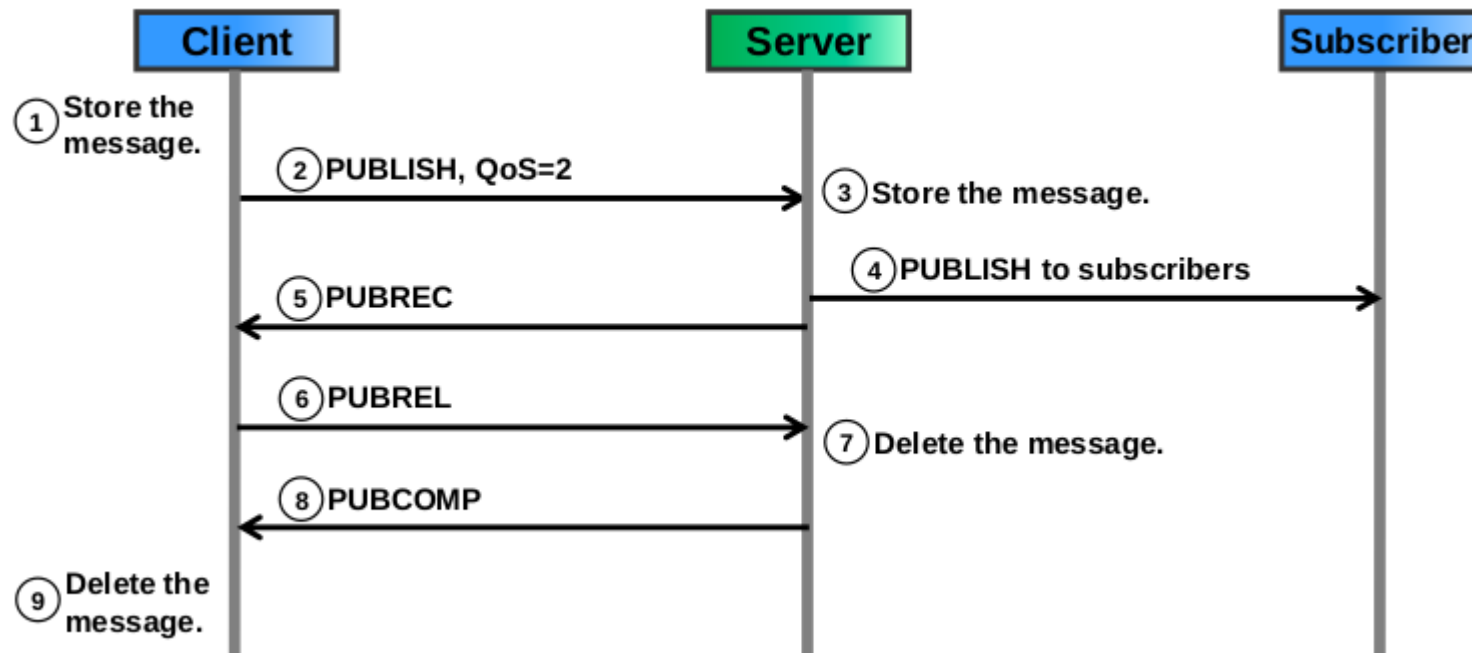
## QoS level 1:

- QoS level 1 affords **at-least-once delivery semantics**. If the client does not receive the **PUBACK** in time, it **re-sends** the message.



## QoS level 2:

- QoS level 2 affords the highest quality delivery semantics **exactly-once**, but comes with the cost of **additional control messages**.





# Thanks!

